

Siddhikesh Arvind Gavit

AI Engineer | Generative AI | Agentic AI | LLM Applications | Gandhinagar, India

+91 93592 79778 | sid.gavit6@gmail.com | [linkedin.com/in/sid-gavit6](https://www.linkedin.com/in/sid-gavit6) | github.com/Siddhu-6 | sid-gavit-portfolio.vercel.app

SUMMARY

AI engineer building agentic AI systems, LLM applications and retrieval-augmented generation (RAG) pipelines in Python. Shipped a LangGraph + MCP multi-agent assistant with human-in-the-loop approvals, a hybrid-search RAG app on Microsoft Azure, and a quantized model on a microcontroller. Focused on production-ready AI: evaluation, guardrails and CI/CD.

EDUCATION

Indian Institute of Information Technology (IIIT) Vadodara

Aug 2023 – May 2027

B.Tech in Computer Science and Engineering

Gandhinagar, Gujarat

Coursework: Machine Learning, Deep Learning, Artificial Intelligence, Data Structures & Algorithms, DBMS, Operating Systems, OOP

Hon. Sharad Pawar International School (CBSE): Class XII 80.00% (2022), Class X 84.80% (2020)

Nashik, Maharashtra

TECHNICAL SKILLS

Generative AI & LLMs: LLM APIs (OpenAI, Anthropic Claude, Google Gemini, Groq), Hugging Face, Ollama, Prompt Engineering, Tool/Function Calling, Structured Outputs, Fine-Tuning & PEFT (LoRA, QLoRA, DoRA), Guardrails, NLP

Agentic AI & RAG: LangGraph, LangChain, CrewAI, LlamaIndex, MCP, Multi-Agent Systems, Human-in-the-Loop, RAG, Hybrid Search, Reranking, Embeddings, Vector Databases (ChromaDB, Pinecone), LLM Evaluation (RAGAS)

ML & Deep Learning: PyTorch, TensorFlow/Keras, scikit-learn, Transformers, XGBoost, LSTM, Pandas, NumPy, Model Quantization (INT8, TensorFlow Lite), Explainable AI (SHAP, LIME)

MLOps & Deployment: FastAPI, REST APIs, Pydantic, Docker, Microsoft Azure (Container Apps), Google Cloud (Cloud Run), GitHub Actions (CI/CD), pytest, Streamlit, Git

Languages, Data & Automation: Python, SQL, C++, Java; PostgreSQL, SQLite; n8n, Zapier, Make.com

EXPERIENCE

Research Intern, Edge AI & Model Optimization

Jun 2026 – Aug 2026

IIIT Vadodara | [GitHub](#)

Gandhinagar, Gujarat

- Compressed a sleep-quality prediction model for on-device use: rebuilt it as a lightweight 1D-CNN and reduced it to approximately 7 KB through INT8 quantization (TensorFlow Lite), with AUROC decreasing only from 0.813 to 0.803.
- Deployed it on an Arduino Nano 33 BLE microcontroller, running each prediction on-device in about 8 ms.
- Co-built the modeling pipeline on simulated wearable data, comparing LSTM, Transformer and XGBoost models with subject-wise splits to avoid data leakage; an LSTM + XGBoost ensemble scored best (AUROC 0.87).

PROJECTS

Synapse: Personal Agentic AI Assistant | *Python, LangGraph, CrewAI, MCP, FastAPI, React, Docker*

[GitHub](#)

- Built a multi-agent assistant that turns natural-language goals into verified actions: LangGraph plans each task, routes steps to specialist agents and CrewAI crews, then checks the results.
- Connected 38 tools through 7 MCP servers (Gmail, Calendar, Notion, Obsidian, n8n, web search and more); risky actions like sending email or deleting notes wait for human-in-the-loop approval.
- Added guardrails against prompt injection, per-run tracing of tokens and latency, automatic fallback from Groq to a local Ollama model, and CI with 70+ tests and 16 behaviour evals.

notes-gpt: Citation-Grounded PDF RAG Assistant | *Python, ChromaDB, Groq, Docker, Azure*

[GitHub](#) · [Live](#)

- Built a RAG app without a framework that answers questions over uploaded PDFs, cites the source page for every claim, and says “I don’t know” when the answer isn’t in the files.
- Combined keyword (BM25) and semantic search over Hugging Face embeddings in ChromaDB, then reranked results with a cross-encoder; answers stream from GPT-OSS 120B via Groq.
- Containerized with Docker and set up CI/CD with GitHub Actions; deployed on Microsoft Azure Container Apps with scale-to-zero, and built a RAGAS harness to evaluate answer quality.

sid.agent: Portfolio AI Agent | *Python, FastAPI, Groq, SSE, Vercel*

[GitHub](#) · [Live](#)

- Built and deployed a public AI agent on the portfolio site that answers recruiter questions about projects and experience, served by a FastAPI backend on Vercel with real-time streaming.
- Added lightweight retrieval over a markdown knowledge base to keep prompts within token limits, plus per-IP rate limiting and input validation to control API usage.

ACHIEVEMENTS & LEADERSHIP

- Hackathons:** Anveshan Hackathon, IIT Guwahati (national level, 48 hours), built an ML model for fake job-posting scam detection; also competed in the Adobe Hackathon and 2x HackIIITV, shipping working prototypes in 24 to 36 hours.
- Leadership:** Paid Teaching Assistant & Football Mentor, IIIT Vadodara (Aug 2025 – Present); Photographer, Obscura Club.
- Football:** 3x state-level player (SGFI National trials); 3rd, Inter-IIIT Sports Meet; 2nd, Concours; 2x runner-up, Ventura.